

Gaussian Mixtures

1.

Modeling Text Documents

2. A Simple Model.

- (a) We shall denote p_{topic} as p , since it is given that this is a single probability. For simplicity, we assume that $y \in \{0, 1\}$, and that $x \in \{0, 1\}^N$. We denote $x[i]$ to be the i th coordinate of the sample x .

Given a sample

$$S = \{(x_1, y_1), \dots, (x_n, y_n)\}$$

We define the following sample statistics. For $x \in \{0, 1\}$, $y \in \{0, 1\}$:

$$n_j(y, x) = |\{i : (x_i, y_i) \in S, x_i[j] = x, y_i = y\}|$$

$$n(y) = |\{i : (x_i, y_i) \in S, y_i = y\}|$$

We want to find estimators for p and for

$$P(x[1] = x_1, \dots, x[N] = x_N | y = y)$$

By the independence of $x[i]|y$, we can simplify this expression:

$$P(x[1] = x_1, \dots, x[N] = x_N | y = y) = \prod_{i=1}^N P(x[i] = x_i | y = y)$$

Thus, we can focus on estimators of p and $P(x[i] = x | y = y) := p_i(x|y)$ (I know that this swaps the arguments of $n_i(y, x)$, it's too much to change now). We should expect our MLEs for p and $p_i(x|y)$ to be the sample means, i.e.

$$\hat{p} = \frac{n(1)}{n}$$

$$\hat{p}_i(x|y) = \frac{n_i(y, x)}{n(y)}$$

We define our log-likelihood function as

$$\ell(\theta|S) = \sum_{i=1}^n \log(P(y = y_i, x[1] = x_i[1], \dots, x[N] = x_i[N]))$$

Given that S was drawn i.i.d., we can simplify.

$$\begin{aligned} \ell(\theta|S) &= \sum_{i=1}^n \log(P(x[1] = x_i[1], \dots, x[N] = x_i[N] | y = y_i) P(y = y_i)) \\ &= \sum_{i=1}^n \log(P(y = y_i) \prod_{j=1}^N P(x[j] = x_i[j] | y = y_i)) \\ &= \sum_{i=1}^n \log(P(y = y_i)) + \sum_{j=1}^N \log(P(x[j] = x_i[j] | y = y_i)) \\ &= \sum_{i=1}^n \log(P(y = y_i)) + \sum_{i=1}^n \sum_{j=1}^N \log(p_j(x_i[j] | y_i)) \end{aligned}$$

Writing the parameters explicitly, we have:

$$\ell(\theta|S) = \sum_{i=1}^n \log(P(y = y_i|p)) + \sum_{i=1}^n \sum_{j=1}^N \log(P(x[i] = x_j[i]|y_i, p_i(x|y)))$$

To solve for the minimum of $\ell(\theta|S)$, we use the method of Lagrange multipliers. First, we can split the problem into two steps. It's clear that that right sum does not depend on p , so we can begin by finding the optimal p .

We note:

$$\begin{aligned} P(y = y_i|p) &= P(y = y_i|y_i = 1, p)P(y_i = 1|p) + P(y = y_i|y_i = 0, p)P(y_i = 0|p) \\ &= P(y = 1|p)[[y_i = 1]] + P(y = 0|p)[[y_i = 0]] \\ &= p^{y_i}(1-p)^{1-y_i} \end{aligned}$$

Plugging this into our log-likelihood, we have:

$$\begin{aligned} \ell(\theta|S) &= \sum_{i=1}^n \log(p^{y_i}(1-p)^{1-y_i}) + \sum_{i=1}^n \sum_{j=1}^N \log(P(x[i] = x_j[i]|y_i, p_i(x|y))) \\ &= \sum_{i=1}^n y_i \log(p) + (1-y_i) \log(1-p) + \sum_{i=1}^n \sum_{j=1}^N \log(P(x[i] = x_j[i]|y_i, p_i(x|y))) \end{aligned}$$

Taking the derivative with respect to p and setting it to zero, we have:

$$\begin{aligned} \frac{d}{dp} \ell(\theta|S) &= \sum_{i=1}^n \frac{y_i}{p} - \frac{1-y_i}{1-p} = 0 \\ \sum_{i=1}^n \frac{y_i}{p} &= \sum_{i=1}^n \frac{1-y_i}{1-p} \\ \frac{1-p}{p} &= \frac{\sum_{i=1}^n 1-y_i}{\sum_{i=1}^n y_i} \\ p &= \frac{\sum_{i=1}^n y_i}{n} \end{aligned}$$

Thus, we have that $\hat{p} = \frac{n(1)}{n}$.

Now, we solve for $\hat{p}_i(x|y)$, by using the method of Lagrange multipliers. Our objective function is as follows:

$$\sum_{i=1}^n \sum_{j=1}^N \log(P(x[j] = x_i[j]|y_i))$$

We can write this in a nicer form.

$$\begin{aligned} \sum_{i=1}^n \sum_{j=1}^N \log(P(x[j] = x_i[j]|y_i)) &= \sum_{i=1}^n \sum_{j=1}^N \log(p_j(x_i[j]|y_i)) \\ &= \sum_{j=1}^N \sum_{i=1}^n \sum_{x \in \{0,1\}} [[x_i[j] = x]] \log(p_j(x|y_i)) \\ &= \sum_{j=1}^N \sum_{i=1}^n \sum_{x \in \{0,1\}} \sum_{y \in \{0,1\}} [[x_i[j] = x \wedge y_i = y]] \log(p_j(x|y)) \\ &= \sum_{j=1}^N \sum_{x \in \{0,1\}} \sum_{y \in \{0,1\}} \log(p_j(x|y)) \sum_{i=1}^n [[x_i[j] = x \wedge y_i = y]] \\ &= \sum_{j=1}^N \sum_{y \in \{0,1\}} \sum_{x \in \{0,1\}} \log(p_j(x|y)) n_j(y, x) \end{aligned}$$

We now have the following constraints:

$$\sum_{x \in \{0,1\}} p_j(x|y) = 1 \quad \forall y \in \{0,1\}, j \in [N]$$

This gives us the following Lagrangian:

$$\mathcal{L} = \sum_{j=1}^N \sum_{y \in \{0,1\}} \sum_{x \in \{0,1\}} \log(p_j(x|y)) n_j(y, x) + \sum_{j=1}^N \sum_{y \in \{0,1\}} \lambda_j(y) \left(\sum_{x \in \{0,1\}} p_j(x|y) - 1 \right)$$

Taking the derivatives with respect to $p_j(x|y)$, we have:

$$\begin{aligned} [p_j(x|y)] : \frac{n_j(y, x)}{p_j(x|y)} &= \lambda_j(y) \\ [\lambda_j(y)] : \sum_{x \in \{0,1\}} p_j(x|y) &= 1 \end{aligned}$$

Since we have equality, in $\lambda_j(y)$, for $x \in \{0,1\}$ we can solve for $p_j(x|y)$:

$$\begin{aligned} \frac{n_j(y, x)}{p_j(x|y)} &= \frac{n_j(y, 1-x)}{p_j(1-x|y)} \\ p_j(1-x|y) &= \frac{n_j(y, 1-x)}{n_j(y, x)} p_j(x|y) \\ \implies 1 &= p_j(x|y) + \frac{n_j(y, 1-x)}{n_j(y, x)} p_j(x|y) \\ n_j(y, x) &= p_j(x|y) n_j(y, x) + n_j(y, 1-x) p_j(x|y) \\ &= p_j(x|y) (n_j(y, x) + n_j(y, 1-x)) \\ p_j(x|y) &= \frac{n_j(y, x)}{n_j(y, x) + n_j(y, 1-x)} \\ \hat{p}_j(x|y) &= \frac{n_j(y, x)}{n(y)} \end{aligned}$$

(b) Using Baye's Law, and conditional independence we have:

$$\begin{aligned} P(Y = 1|X = x) &= \frac{P(X = x|Y = 1)P(Y = 1)}{P(X = x)} \\ &= \frac{P(X[1] = x[1], \dots, X[N] = x[N]|Y = 1)P(Y = 1)}{P(X[1] = x[1], \dots, X[N] = x[N])} \\ &= \frac{P(Y = 1) \prod_{i=1}^N P(X[i] = x[i]|Y = 1)}{P(X[1] = x[1], \dots, X[N] = x[N]|Y = 1)P(Y = 1) + P(X[1] = x[1], \dots, X[N] = x[N]|Y = 0)P(Y = 0)} \\ &= \frac{p \prod_{i=1}^N p_i(x[i]|1)}{p \prod_{i=1}^N p_i(x[i]|1) + (1-p) \prod_{i=1}^N p_i(x[i]|0)} \end{aligned}$$

Now we can reduce this into the form of a logistic function.

$$\begin{aligned} P(Y = 1|X = x) &= \frac{p \prod_{i=1}^N p_i(x[i]|1)}{p \prod_{i=1}^N p_i(x[i]|1) + (1-p) \prod_{i=1}^N p_i(x[i]|0)} \\ &= \frac{1}{1 + \frac{1-p}{p} \frac{\prod_{i=1}^N p_i(x[i]|0)}{\prod_{i=1}^N p_i(x[i]|1)}} \\ &= \frac{1}{1 + e^{-(\log(\frac{p}{1-p}) + \sum_{i=1}^N \log(\frac{p_i(x[i]|1)}{p_i(x[i]|0)}))}} \\ &\quad 3 \end{aligned}$$

Therefore, we can get our discriminant as follows:

$$r(x) = \log \left(\frac{p}{1-p} \right) + \sum_{i=1}^N \log \left(\frac{p_i(x[i]|1)}{p_i(x[i]|0)} \right)$$

(c) We can simplify the discriminant by noting

$$p_i(x|y) = p_i(1|y)^x p_i(0|y)^{1-x}$$

Giving us

$$\begin{aligned} r(x) &= \log \left(\frac{p}{1-p} \right) + \sum_{i=1}^N \log \left(\frac{p_i(1|1)^{x[i]} p_i(0|1)^{1-x[i]}}{p_i(1|0)^{x[i]} p_i(0|0)^{1-x[i]}} \right) \\ &= \log \left(\frac{p}{1-p} \right) + \sum_{i=1}^N \left(x[i] \log \left(\frac{p_i(1|1)}{p_i(1|0)} \right) + (1-x[i]) \log \left(\frac{p_i(0|1)}{p_i(0|0)} \right) \right) \\ &= \log \left(\frac{p}{1-p} \right) + \sum_{i=1}^N x[i] \log \left(\frac{p_i(1|1)}{p_i(1|0)} \right) - \sum_{i=1}^N x[i] \log \left(\frac{p_i(0|1)}{p_i(0|0)} \right) + \sum_{i=1}^N \log \left(\frac{p_i(0|1)}{p_i(0|0)} \right) \\ &= \log \left(\frac{p}{1-p} \right) + \sum_{i=1}^N x[i] \left(\log \left(\frac{p_i(1|1)}{p_i(1|0)} \right) - \log \left(\frac{p_i(0|1)}{p_i(0|0)} \right) \right) + \sum_{i=1}^N \log \left(\frac{p_i(0|1)}{p_i(0|0)} \right) \\ &= \log \left(\frac{p}{1-p} \right) + \sum_{i=1}^N \log \left(\frac{p_i(0|1)}{p_i(0|0)} \right) + \sum_{i=1}^N x[i] \left(\log \left(\frac{p_i(1|1)}{p_i(0|1)} \frac{p_i(0|0)}{p_i(1|0)} \right) \right) \end{aligned}$$

The feature map must include a constant 1 to account for the term on the left, and must have N more features for each of $x[i]$. Thus, our feature map is simply:

$$\phi : x \mapsto (1, x[1], \dots, x[N])$$

Therefore, our vector w , such that $r(x) = \langle w, \phi(x) \rangle$, is:

$$w = \left(\log \left(\frac{p}{1-p} \right) + \sum_{i=1}^N \log \left(\frac{p_i(0|1)}{p_i(0|0)} \right), \log \left(\frac{p_1(1|1)}{p_1(0|1)} \frac{p_1(0|0)}{p_1(1|0)} \right), \dots, \log \left(\frac{p_N(1|1)}{p_N(0|1)} \frac{p_N(0|0)}{p_N(1|0)} \right) \right)$$

(d) The log odds term in the bias has a simple interpretation.

$$\begin{aligned} \frac{\hat{p}}{1-\hat{p}} &= \frac{n(1)/n}{n(0)/n} = \frac{n(1)}{n(0)} \\ \log \left(\frac{\hat{p}}{1-\hat{p}} \right) &= \log \left(\frac{n(1)}{n(0)} \right) \end{aligned}$$

Similarly,

$$\frac{\hat{p}_i(x|y)}{\hat{p}_i(x'|y')} = \frac{n_i(y, x)/n(y)}{n_i(y', x')/n(y')}$$

So,

$$\begin{aligned} \frac{\hat{p}_i(0|1)}{\hat{p}_i(0|0)} &= \frac{n_i(1, 0)/n(1)}{n_i(0, 0)/n(0)} \\ \frac{\hat{p}_i(1|1)\hat{p}_i(0|0)}{\hat{p}_i(0|1)\hat{p}_i(1|0)} &= \frac{n_i(1, 1)/n(1)n_i(0, 0)/n(0)}{n_i(1, 0)/n(1)n_i(0, 1)/n(0)} \end{aligned}$$

$$= \frac{n_i(1,1)n_i(0,0)}{n_i(1,0)n_i(0,1)}$$

So we have the following simplification for w :

$$w = \left((N-1) \log \frac{n(0)}{n(1)} + \sum_{i=1}^N \log \frac{n_i(1,0)}{n_i(0,0)}, \log \frac{n_1(1,1)n_1(0,0)}{n_1(1,0)n_1(0,1)}, \dots, \log \frac{n_N(1,1)n_N(0,0)}{n_N(1,0)n_N(0,1)} \right)$$

3. Adding a Prior.

(a) The MAP estimate is defined as follows:

$$\hat{\theta} = \arg \max_{\theta} p(\theta|S)$$

In our case,

$$\theta = (p, \{p_y\})$$

Where we define:

$$p = P(y = 1)$$

$$p_y[i] = P(x[i] = 1|y)$$

and p_y is a vector of N elements. Let S be a sample of n i.i.d. points.

$$S = ((x_1, y_1), \dots, (x_n, y_n))$$

our posterior distribution, $p(\theta|S)$ is given by:

$$\begin{aligned} p(p, \{p_y\}|S) &= \frac{p(S|p, \{p_y\})p(p, \{p_y\})}{p(S)} \\ &= \frac{p(X|Y, p, \{p_y\})p(Y|p, \{p_y\})p(p, \{p_y\})}{p(S)} \\ &= \frac{p(X|Y, \{p_y\})p(Y|p)p(p, \{p_y\})}{p(X|Y)p(Y)} \end{aligned}$$

where X is the vector of x_i 's and Y is the vector of y_i 's.

Note that, we are not conditioning the denominator with respect to the parameters we are optimizing over. The denominator is the distribution over the distributions of p and $\{p_y\}$. Therefore, we can ignore it in the optimization problem.

$$\hat{\theta} = \arg \max_{p, \{p_y\}} p(X|Y, \{p_y\})p(Y|p)p(p, \{p_y\})$$

We break this expression down, term by term, first focusing on the last term.

$$\begin{aligned} p(p, \{p_y\}) &= p(p)p(\{p_y\}) = f_{Dir(1)}(p)p(p_1)p(p_0) \\ &= f_{Dir(\alpha)}(p_1)f_{Dir(\alpha)}(p_0) \\ &= \frac{1}{Z(\alpha)^2} \prod_{i=1}^N p_1[i]^{\alpha-1} p_0[i]^{\alpha-1} \end{aligned}$$

Since $Z(\alpha)^2$ is fixed, we can ignore it in the expression for $\hat{\theta}$. Now we focus on the second term.

$$\begin{aligned} p(Y|p) &= P(Y_1 = y_1, \dots, Y_n = y_n|p) = \prod_{i=1}^n P(Y_i = y_i|p) \\ &= \prod_{i=1}^n p^{y_i} (1-p)^{1-y_i} \end{aligned}$$

Now we focus on the first term.

$$p(X|Y, \{p_y\}) = P(X_1 = x_1, \dots, X_n = x_n|Y_1 = y_1, \dots, Y_n = y_n, \{p_y\})$$

$$\begin{aligned}
&= \prod_{i=1}^n P(X_i = x_i | Y_1 = y_1, \dots, Y_n = y_n, \{p_y\}) \\
&= \prod_{i=1}^n P(X_i = x_i | Y_i = y_i, \{p_y\}) \\
&= \prod_{i=1}^n P(X_i[1] = x_i[1], \dots, X_i[N] = x_i[N] | Y_i = y_i, \{p_y\}) \\
&= \prod_{i=1}^n \prod_{j=1}^N P(X_i[j] = x_i[j] | Y_i = y_i, \{p_y\})
\end{aligned}$$

Since log is monotone, we can take the log of our expression to get the arg max.

$$\begin{aligned}
\hat{\theta} &= \arg \max_{p, \{p_y\}} \sum_{i=1}^n \sum_{j=1}^N \log P(X_i[j] = x_i[j] | Y_i = y_i, \{p_y\}) \\
&\quad + \sum_{i=1}^n y_i \log(p) + (1 - y_i) \log(1 - p) \\
&\quad + \sum_{i=1}^N \sum_{y \in \{0,1\}} (\alpha - 1) \log(p_y[i])
\end{aligned}$$

First, we get \hat{p} by differentiating with respect to p and setting it to zero.

$$\begin{aligned}
\frac{d}{dp} \hat{\theta} &= \sum_{i=1}^n \frac{y_i}{p} - \frac{1 - y_i}{1 - p} = 0 \\
\sum_{i=1}^n \frac{y_i}{p} &= \sum_{i=1}^n \frac{1 - y_i}{1 - p} \\
\frac{1 - p}{p} &= \frac{\sum_{i=1}^n 1 - y_i}{\sum_{i=1}^n y_i} \\
p &= \frac{\sum_{i=1}^n y_i}{n} = \frac{n(1)}{n}
\end{aligned}$$

Where $n(y)$ is the number of y_i 's that are equal to y . Before we try and solve for $p_y[i]$, we can do a better job at simplifying the first term.

$$\begin{aligned}
\log P(X_i[j] = x_i[j] | Y_i = y_i, \{p_y\}) &= [[x_i[j] = 1]] \log(p_{y_i[j]}) + [[x_i[j] = 0]] \log(1 - p_{y_i[j]}) \\
&= \log(p_{y_i[j]}^{x_i[j]} (1 - p_{y_i[j]})^{1 - x_i[j]}) \\
&= \sum_{y \in \{0,1\}} [[y_i = y]] \log(p_y[j]^{x_i[j]} (1 - p_y[j])^{1 - x_i[j]}) \\
\Rightarrow \sum_{i=1}^n \sum_{j=1}^N \log P(X_i[j] = x_i[j] | Y_i = y_i, \{p_y\}) &= \sum_{i=1}^n \sum_{j=1}^N \sum_{y \in \{0,1\}} [[y_i = y]] \log(p_y[j]^{x_i[j]} (1 - p_y[j])^{1 - x_i[j]}) \\
&= \sum_{j=1}^N \sum_{i=1}^n \sum_{y \in \{0,1\}} \sum_{x \in \{0,1\}} [[y_i = y \wedge x_i[j] = x]] \log(p_y[j]^x (1 - p_y[j])^{1 - x})
\end{aligned}$$

$$\begin{aligned}
&= \sum_{j=1}^N \sum_{y \in \{0,1\}} \sum_{x \in \{0,1\}} \log(p_y[j]^x (1 - p_y[j])^{1-x}) \sum_{i=1}^n [[y_i = y \wedge x_i[j] = x]] \\
&= \sum_{j=1}^N \sum_{y \in \{0,1\}} \sum_{x \in \{0,1\}} \log(p_y[j]^x (1 - p_y[j])^{1-x}) n_j(x, y) \\
&= \sum_{j=1}^N \sum_{y \in \{0,1\}} \sum_{x \in \{0,1\}} n_j(x, y) (x \log(p_y[j]) + (1 - x) \log(1 - p_y[j]))
\end{aligned}$$

Where $n_j(x, y)$ is the number of (x_i, y_i) 's such that $x_i[j] = x$ and $y_i = y$. We differentiate with respect to $p_y[j]$ and set equal to zero.

$$\begin{aligned}
\frac{d}{dp_y[j]} \hat{\theta} &= \sum_{x \in \{0,1\}} n_j(x, y) \left(\frac{x}{p_y[j]} - \frac{1-x}{1-p_y[j]} \right) + (\alpha - 1) \frac{1}{p_y[j]} \\
&= n_j(1, y) \frac{1}{p_y[j]} - n_j(0, y) \frac{1}{1-p_y[j]} + (\alpha - 1) \frac{1}{p_y[j]} = 0 \\
\Rightarrow \frac{1}{1-p_y[j]} n_j(0, y) &= \frac{1}{p_y[j]} (n_j(1, y) + (\alpha - 1)) \\
\Rightarrow \frac{1-p_y[j]}{p_y[j]} &= \frac{n_j(0, y)}{n_j(1, y) + (\alpha - 1)} \\
\Rightarrow p_y[j] &= \frac{n_j(1, y) + (\alpha - 1)}{n_j(0, y) + n_j(1, y) + (\alpha - 1)} \\
\Rightarrow \hat{p}_y[j] &= \frac{n_j(1, y) + (\alpha - 1)}{n(y) + (\alpha - 1)}
\end{aligned}$$

(b) Recall that the discriminant is given by

$$r(x) = \log \frac{p}{1-p} + \sum_{i=1}^N \log \frac{p_1[i]^{x[i]} (1 - p_1[i])^{1-x[i]}}{p_0[i]^{x[i]} (1 - p_0[i])^{1-x[i]}}$$

We can simplify this

$$\begin{aligned}
r(x) &= \log \frac{p}{1-p} + \sum_{i=1}^N \sum_{x \in \{0,1\}} [[x[i] = x]] \left(x \log \frac{p_1[i]}{p_0[i]} + (1-x) \log \frac{1-p_1[i]}{1-p_0[i]} \right) \\
&= \log \frac{p}{1-p} + \sum_{i=1}^N [[x[i] = 1]] \log \frac{p_1[i]}{p_0[i]} + [[x[i] = 0]] \log \frac{1-p_1[i]}{1-p_0[i]}
\end{aligned}$$

Taking a slightly modified feature map from before, i.e.

$$\phi : x \mapsto (x[1], \dots, x[N])$$

We have the following $w \in \mathbb{R}^N$ and $b \in \mathbb{R}$ such that $r(x) = \langle w, \phi(x) \rangle + b$:

$$\begin{aligned}
w[i] &= \log \frac{p_1[i]}{p_0[i]} - \log \frac{1-p_1[i]}{1-p_0[i]} \\
b &= \log \frac{p}{1-p} + \sum_{i=1}^N \log \frac{1-p_1[i]}{1-p_0[i]}
\end{aligned}$$

Now we can easily plug in our MAP estimators.

$$w[i] = \log \frac{n_i(1, 1) + (\alpha - 1)}{n(1) + (\alpha - 1)} \frac{n(0) + (\alpha - 1)}{n_i(1, 0) + (\alpha - 1)} - \log \frac{n_i(0, 1)}{n(1) + (\alpha - 1)} \frac{n(0) + (\alpha - 1)}{n_i(0, 0)}$$

$$\begin{aligned}
&= \log \frac{n_i(1, 1) + (\alpha - 1)}{n_i(1, 0) + (\alpha - 1)} - \log \frac{n_i(0, 1)}{n_i(0, 0)} \\
b &= \log \frac{n(1)/n}{n(0)/n} + \sum_{i=1}^N \log \frac{n_i(0, 1)}{n(1) + (\alpha - 1)} \frac{n(0) + (\alpha - 1)}{n_i(0, 0)} \\
&= \log \frac{n(1)}{n(0)} + \sum_{i=1}^N \log \frac{n_i(0, 1)}{n(1) + (\alpha - 1)} \frac{n(0) + (\alpha - 1)}{n_i(0, 0)}
\end{aligned}$$

4. Multiple Classes.

5. Adding Dependencies: A Markov Model.